

Bootstrapping planning action models

Guillem Francés

Information and Communication Technology
Universitat Pompeu Fabra
Roc Boronat 138, 08018 Barcelona, Spain
@upf.edu

Sergio Jiménez

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
sjimenez@unimelb.edu.au

Nir Lipovetzky

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
@unimelb.edu.au

Miguel Ramírez

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
@unimelb.edu.au

Abstract

This paper presents an innovative approach for learning classical planning action models from minimal input knowledge and using exclusively existing classical planners. First, the paper defines a classical planning compilation to learn action models from example plans that are labeled with their corresponding initial and final states. Second, the paper explains how to collect informative plan examples using a classical planner based on pure exploratory search.

Introduction

Off-the-shelf planners reason about action models that correctly and completely capture the possible world transitions (Geffner and Bonet 2013). Building such models is complex even for planning experts (Kambhampati 2007).

In Machine Learning (ML) models are not hand-coded but computed from examples (Michalski, Carbonell, and Mitchell 2013). Unfortunately, the application of off-the-shelf ML techniques to learning planning action models is not straightforward. On the one hand the collection of *informative* examples for learning planning action models is complex. Planning actions traditionally include preconditions that are only satisfied by specific sequences of actions and often with a low probability of being chosen by chance. Therefore simple exploration approaches, s.t. random walks, easily under-sample planning state spaces (Fern, Yoon, and Givan 2004). On the other hand the traditional output of off-the-shelf ML techniques is a scalar value (an integer, in the case of classification tasks, or a real value, in the case of regression tasks).

In this work we focus on learning action models for classical planning. This is a well-studied problem where the dynamics of a given action can be captured lifting the literals that change between the pre and post-state of an action execution. There are sophisticated learning approaches for this task, like ARMS (Yang, Wu, and Jiang 2007) or LOCM (Cresswell, McCluskey, and West 2013) systems that do not require full knowledge of the states traversed by the example plans.

Despite these systems' achievements they still assume that example plans are given from an external agent. Mo-

tivated by recent advances on effective exploration of planning state spaces () and on the learning of complex structures with planners (Segovia-Aguas, Jiménez, and Jonsson 2017), this paper introduces an innovative approach for learning classical planning action models. The contribution of this work is then two-fold:

1. An inductive learning algorithm that minimizes the burden of required supervision and that can be defined as a classical planning compilation.
2. A method for autonomously collect *informative* the example plans for action model learning using an exploration-based classical planner.

Background

Here we define the planning models we use on this work.

Classical Planning

We use F to denote the set of *fluents* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$, i.e. $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (WLOG we assume that L does not assign conflicting values to any fluent). We use $\mathcal{L}(F)$ to denote the set of all literal sets on F , i.e. all partial assignments of values to fluents.

A *state* s is a total assignment of values to fluents, i.e. $|s| = |F|$ and the number of states is $2^{|F|}$. Explicitly including negative literals $\neg f$ in states simplifies subsequent definitions, but we often abuse notation by defining a state s only in terms of the fluents that are true in s , as is common in STRIPS planning.

Under this formalism, a *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of actions. Each action $a \in A$ has a set of literals $\text{pre}(a) \in \mathcal{L}(F)$ called the *precondition*, a set of positive effects $\text{add}(a) \in \mathcal{L}(F)$ and a set of del effects $\text{del}(a) \in \mathcal{L}(F)$. An action $a \in A$ is applicable in state s iff $\text{pre}(a) \subseteq s$, and the result of applying a in s is a new state $\theta(s, a) = (s \setminus \neg \text{del}(a)) \cup \text{add}(a)$.

Given a planning frame $\Phi = \langle F, A \rangle$, a *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is an initial state and G is a goal condition, i.e. a set of literals on F . A *plan* for P is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$ that induces a state sequence $\langle s_0, s_1, \dots, s_n \rangle$ such that $s_0 = I$

and, for each i such that $1 \leq i \leq n$, a_i is applicable in s_{i-1} and generates the successor state $s_i = \theta(s_{i-1}, a_i)$. The plan π solves P if and only if $G \subseteq s_n$, i.e. if the goal condition is satisfied following the application of π in I .

We assume that fluents are instantiated from predicates, as in PDDL (Fox and Long 2003). Specifically, there exists a set of predicates Ψ , and each predicate $p \in \Psi$ has an argument list of arity $ar(p)$. Given a set of objects Ω , the set of fluents F is then induced by assigning objects in Ω to the arguments of predicates in Ψ , i.e. $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{ar(p)}\}$ where, given a set X , X^n is the n -th Cartesian power of X .

Likewise we assume that actions in $a \in A$ are instantiated from operator schema Ξ , i.e. $A = \{\xi(\omega) : \xi \in \Xi, \omega \in \Omega^{ar(\xi)}\}$. An operator schema $\xi \in \Xi$, is represented by an operator *header*: that contains a unique symbol, $name(\xi)$, and a list of variables, $vars(\xi)$. The operator *body* consists of three sets of lifted predicates: the *preconditions*, $pre(\xi)$, the *positive effects*, $add(\xi)$, and the *negative effects*, $del(\xi)$. Figure 1 shows an example of planning action schema from the blockworld as represented in PDDL where positive and negative effects are grouped together.

```
(:action stack
:parameters (?x1 ?x2)
:precondition (and (holding ?x1)
                  (clear ?x2))
:effect (and (not (holding ?x1))
             (not (clear ?x2))
             (clear ?x1)
             (handempty)
             (on ?x1 ?x2)))
```

Figure 1: Example of the *stack* planning action schema from the blockworld as represented in PDDL.

Classical Planning with Conditional Effects

Conditional effects make it possible to repeatedly refer to the same action even though their precise effects depend on the current state. In this case each action $a \in A$ has a set of literals $pre(a) \in \mathcal{L}(F)$ called the *precondition* and a set of conditional effects $cond(a)$. Each conditional effect $C \triangleright E \in cond(a)$ is composed of two sets of literals $C \in \mathcal{L}(F)$ (the condition) and $E \in \mathcal{L}(F)$ (the effect). An action $a \in A$ is applicable in state s if and only if $pre(a) \subseteq s$, and the resulting set of *triggered effects* is

$$eff(s, a) = \bigcup_{C \triangleright E \in cond(a), C \subseteq s} E,$$

i.e. effects whose conditions hold in s . The result of applying a in s is a new state $\theta(s, a) = (s \setminus \neg eff(s, a)) \cup eff(s, a)$.

Learning planning action models

This section formalizes the learning task of computing a planning action model from a given set of lifted predicates and labelled input plans. This task is defined as $\langle \Psi, \Pi, \Lambda \rangle$:

- Ψ is the set of lifted predicates that define the abstract state space of a given planning domain,

- $\Pi = \{\pi_1, \dots, \pi_t\}$ is the given set of example plans,
- $\Lambda = \{\lambda_1, \dots, \lambda_t\}$ is the corresponding set of labels s.t., each plan π_i , $1 \leq i \leq t$, has an associated label $\lambda_i = (s_i, s'_i)$ such that s'_i is the state resulting from executing π_i starting from the state s_i .

A solution to the $\langle \Psi, \Pi, \Lambda \rangle$ learning task is a set of operator schema Ξ that is compliant with the predicates in Ψ , the example plans Π , and their labels Λ .

Learning action models using a classical planner

Our approach is compiling the learning task $\langle \Psi, \Pi, \Lambda \rangle$ into a classical planning task $P' = \langle F', A', I', G' \rangle$ that can later be solved by an off-the-shelf classical planner. The intuition behind the compilation is that a solution to P' is a sequence of actions that first, programs the action action model (the preconditions, del and add effects of each action schema $\xi \in \Xi$) and then, sequentially validates the programmed action model in the example plans Π (and their labels Λ), one after the other.

To formalize the compilation we first define t classical planning instances $P_1 = \langle F, A, I_1, G_1 \rangle, \dots, P_t = \langle F, A, I_t, G_t \rangle$, that belong to the same planning frame $\Phi = \langle F, A \rangle$ (i.e. share the same fluents and actions and differ only in the initial state and goals). Let Ω be the set of objects that appear either in the example plans Π or in their labels Λ , i.e., $\Omega = \{o | o \in \pi_i \vee o \in \lambda_i, 1 \leq i \leq t\}$. Then F and A are the set of fluents and actions built instantiating the predicates in Ψ with the objects in Ω . Finally the initial state I_i , $1 \leq i \leq t$, is given by the state $s_i \in \lambda_i$ and the goals are defined by the state $s'_i \in \lambda_i$. According to these definitions each plan $\pi_i \in \Pi$, $1 \leq i \leq t$, is a solution to the corresponding classical planning instance $P_i = \langle F, A, I_i, G_i \rangle$.

Now we are ready to define the compilation for learning action models using a classical planner. Given a learning task $\langle \Psi, \Pi, \Lambda \rangle$ the compilation outputs a classical planning task $P' = \langle F', A', I', G' \rangle$ where:

- F' extends F with:
 - Fluents representing the programmed action model. They have the form $header(name(\xi), \Omega_v^{ar(\xi)})$, $pre_p(name(\xi), \Omega_v^{ar(p)})$, $del_p(name(\xi), \Omega_v^{ar(p)})$ and $add_p(name(\xi), \Omega_v^{ar(p)})$ where $p \in \Psi$, $\xi \in \Xi$ and $\Omega_v = \{var_1, \dots, var_v\}$ is a new set of objects, $\Omega_v \cap \Omega = \emptyset$, representing variable names. The number of *variable* objects, i.e. $|\Omega_v|$, is given by the action with the maximum arity in Π . For instance, for the blockworld $\Omega_v = \{var_1, var_2\}$ since actions *stack* and *unstack* have two parameters.
 - Fluents $\{test_i\}_{1 \leq i \leq t}$, indicating the plan where the programmed model is currently being validated and fluents $plan(name(\xi), i, \Omega^{ar(a)})$ for encoding those plans. Fluents at_i and $next_{i,i2}$, $1 \leq i < i2 \leq n$, indicating the plan step where the programmed model is being validated and n the max length of a plan in Π .
- I' , contains the fluents from F that encode the initial state $s_1 \in P_1$, the fluents $plan(name(\xi), i, \Omega^{ar(a)})$, $1 \leq i \leq |\pi_1|$ that encode the plan $\pi_1 \in \Pi$ for solving P_1 , and the

fluents at_1 and $\{next_{i,i2}\}$, $1 \leq i < i2 \leq n$, for indicating that the plan step where to start validating the programmed model.

- $G' = \{test_i\}$, $1 \leq i \leq t$, indicates that the programmed model is validated in all the plans in Π .
- A' replaces the actions in A with actions of three types:

1. The actions for programming an action schema:

- A *precondition* with predicate $p \in \Psi$ and variables $v \in \Omega_v^{ar(p)}$ in the action schema $\xi \in \Xi$:

$$\begin{aligned} \text{pre}(\text{programPre}_{\xi,p(v)}) &= \{\neg \text{pre}_{\xi}(p(v)), \neg \text{add}_{\xi}(p(v))\}. \\ \text{cond}(\text{programPre}_{\xi,p(v)}) &= \{\emptyset\} \triangleright \{\text{pre}_{\xi}(p(v))\}. \end{aligned}$$

- A *negative effect* with predicate $p \in \Psi$ and variables $v \in \Omega_v^{ar(p)}$ in the action schema $\xi \in \Xi$:

$$\begin{aligned} \text{pre}(\text{programDel}_{\xi,p(v)}) &= \{\text{pre}_{\xi}(p(v)), \neg \text{del}_{\xi}(p(v)), \\ &\quad \neg \text{add}_{\xi}(p(v))\}. \\ \text{cond}(\text{programDel}_{\xi,p(v)}) &= \{\emptyset\} \triangleright \{\text{del}_{\xi}(p(v))\}. \end{aligned}$$

- A *positive effect* with predicate $p \in \Psi$ and variables $v \in \Omega_v^{ar(p)}$ in the action schema $\xi \in \Xi$:

$$\begin{aligned} \text{pre}(\text{programAdd}_{\xi,p(v)}) &= \{\neg \text{pre}_{\xi}(p(v)), \neg \text{del}_{\xi}(p(v)), \\ &\quad \neg \text{add}_{\xi}(p(v))\}. \\ \text{cond}(\text{programAdd}_{\xi,p(v)}) &= \{\emptyset\} \triangleright \{\text{add}_{\xi}(p(v))\}. \end{aligned}$$

2. The actions for applying an operator schema $\xi \in \Xi$ (that is already programmed with variables $v \in \Omega_v^{ar(\xi)}$) and that is bound with objects $v' \in \Omega^{ar(\xi)}$

$$\begin{aligned} \text{pre}(\text{apply}_{\xi,v,v'}) &= \{\neg \text{header}(\xi(v))\} \cup \\ &\quad \{\neg \text{pre}_{\xi}(p(v)) \vee p(v')\}_{\forall p \in \Psi}. \\ \text{cond}(\text{apply}_{\xi,v,v'}) &= \{\emptyset\} \triangleright \{\text{header}(\xi(v))\}, \\ &\quad \{at_i\} \triangleright \{\neg at_i, at_{i+1}\}_{\forall i \in 1 \leq i < n}, \\ &\quad \{\text{del}_{\xi}(p(v))\} \triangleright \{\neg p(v')\}_{\forall p \in \Psi}, \\ &\quad \{\text{add}_{\xi}(p(v))\} \triangleright \{p(v')\}_{\forall p \in \Psi}. \end{aligned}$$

3. The actions for changing the active test, that is the plan where the model is currently being validated.

$$\begin{aligned} \text{pre}(\text{validate}_i) &= G_i \cup \{at_{|\Pi_i|}\}, \{test_j\}_{j \in 1 \leq j < i\}. \\ \text{cond}(\text{validate}_i) &= \{\emptyset\} \triangleright \{test_i\}, \\ &\quad \{\emptyset\} \triangleright \{\neg at_{|\Pi_i|}, at_1\}, \\ &\quad \{\emptyset\} \triangleright \{\neg \text{plan}(\xi, k, v)\}_{1 \leq k < |\Pi_i|}, \\ &\quad \{\emptyset\} \triangleright \{\text{plan}(\xi, k, v)\}_{1 \leq k < |\Pi_{i+1}|}. \end{aligned}$$

Lemma 1. Any classical plan π that solves $P' = \langle F', A', I', G' \rangle$ induces a valid action model that solves the learning task $\langle \Psi, \Pi, \Lambda \rangle$.

Proof sketch. Once an action schema is programmed it can only be executed. The only way of achieving a test fluent is by executing the actions indicated by the corresponding plan and achieve the goal state in its associated label. If this is done for all the input examples it means that the programmed model is compliant with the learning input knowledge and hence, solves the action model learning task. \square

Interestingly the compilation is also valid for partially specified action models since the known preconditions and effects can be part of the initial state. With this regard the approach allows also transfer learning in which we generate the action model for a given sub-task and then encode this model as already programmed actions for learning new or more challenging action models.

Generating the learning examples

Observation: If there is no novel predicates no learning is possible. Possible Solution: Novelty over lifted predicates?

Evaluation

How to compare two action models? Possible Solution: over test cases

Related work

Conclusions

References

- Cresswell, S. N.; McCluskey, T. L.; and West, M. M. 2013. Acquiring planning domain models using locm. *The Knowledge Engineering Review* 28(02):195–213.
- Fern, A.; Yoon, S. W.; and Givan, R. 2004. Learning domain-specific control knowledge from random walks. In *ICAPS*, 191–199.
- Fox, M., and Long, D. 2003. Pddl2. 1: An extension to pddl for expressing temporal planning domains. *J. Artif. Intell. Res. (JAIR)* 20:61–124.
- Geffner, H., and Bonet, B. 2013. A concise introduction to models and methods for automated planning.
- Kambhampati, S. 2007. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *Proceedings of the National Conference on Artificial Intelligence*, volume 22, 1601. Menlo Park, CA; Cambridge, MA; London; AAAI Press; MIT Press; 1999.
- Michalski, R. S.; Carbonell, J. G.; and Mitchell, T. M. 2013. *Machine learning: An artificial intelligence approach*. Springer Science & Business Media.
- Segovia-Aguas, J.; Jiménez, S.; and Jonsson, A. 2017. Un-supervised learning of planning tasks. In *International Conference on Automated Planning and Scheduling*.
- Yang, Q.; Wu, K.; and Jiang, Y. 2007. Learning action models from plan examples using weighted max-sat. *Artificial Intelligence* 171(2-3):107–143.