

Bootstrapping classical planning action models

Guillem Francés

Information and Communication Technologies
Universitat Pompeu Fabra
Roc Boronat 138, 08018 Barcelona, Spain
@upf.edu

Sergio Jiménez

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
sjimenez@unimelb.edu.au

Nir Lipovetzky

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
@unimelb.edu.au

Miguel Ramírez

Computing and Information Systems
University of Melbourne
Parkville, Victoria 3010, Australia
@unimelb.edu.au

Abstract

This paper presents a novel approach for learning classical planning action models from minimal input knowledge and using exclusively existing classical planners. First, the paper defines a classical planning compilation to learn action models from examples. Interestingly this compilation accepts partially specified action models and allows different levels of available input knowledge in the examples. Second, the paper explains how to autonomously collect informative examples using a classical planner based on pure exploratory search.

Introduction

Off-the-shelf planners reason about action models that correctly and completely capture the possible world transitions (Geffner and Bonet 2013). Unfortunately building such models is complex even for planning experts limiting the potential of automated planning (Kambhampati 2007).

Machine Learning (ML) is able to compute models from examples (Michalski, Carbonell, and Mitchell 2013) but applying off-the-shelf ML to learning planning action models is not straightforward. First, the inputs of ML algorithms are usually a finite set of values encoding some features of objects. Planning tasks include actions and are closer to the representation of procedures and behaviours than to object representation. Second, the traditional output of off-the-shelf ML techniques is a scalar value (an integer, in the case of classification tasks, or a real value, in the case of regression tasks). The output for learning classical planning actions is not a scalar but a declarative generative model of the possible state transitions. Last but not least, the collection of *informative* examples for learning planning action models is challenging. Planning actions include preconditions that are only satisfied by specific sequences of actions and often with a low probability of being chosen by chance (Fern, Yoon, and Givan 2004).

Learning classical planning action models is a well-studied problem with sophisticated algorithms, like ARMS (Yang, Wu, and Jiang 2007) or LOCM (Cresswell, McCluskey, and West 2013) that do not require the full knowledge of all the states traversed by an example plan. Motivated by recent advances on learning generative models

using classical planners (Segovia-Aguas, Jiménez, and Jonsson 2017) and on the effective exploration of planning state spaces (Francs et al. 2017), this paper aims introducing an innovative approach for learning classical planning action models. The contribution of the work is two-fold:

1. An inductive learning algorithm that minimizes the required input knowledge, that is flexible to different levels of available input knowledge, and that can be defined as a classical planning compilation.
2. A method for autonomously collect *informative* examples for action model learning using an exploration-based classical planner.

Background

This section defines the planning models used on this work.

Classical Planning

We use F to denote the set of *fluents* (propositional variables) describing a state. A *literal* l is a valuation of a fluent $f \in F$, i.e. $l = f$ or $l = \neg f$. A set of literals L represents a partial assignment of values to fluents (WLOG we assume that L does not assign conflicting values to any fluent). We use $\mathcal{L}(F)$ to denote the set of all literal sets on F , i.e. all partial assignments of values to fluents. A *state* s is a total assignment of values to fluents, i.e. $|s| = |F|$, so the number of states is $2^{|F|}$.

A *classical planning frame* is a tuple $\Phi = \langle F, A \rangle$, where F is a set of fluents and A is a set of actions. Each action $a \in A$ has a set of literals $\text{pre}(a) \in \mathcal{L}(F)$ called the *preconditions*, a set of positive effects $\text{add}(a) \in \mathcal{L}(F)$ and a set of negative effects $\text{del}(a) \in \mathcal{L}(F)$. An action $a \in A$ is applicable in state s iff $\text{pre}(a) \subseteq s$, and the result of applying a in s is a new state $\theta(s, a) = (s \setminus \neg \text{del}(a)) \cup \text{add}(a)$.

A *classical planning problem* is a tuple $P = \langle F, A, I, G \rangle$, where I is an initial state and $G \in \mathcal{L}(F)$ is a goal condition. A *plan* for P is an action sequence $\pi = \langle a_1, \dots, a_n \rangle$ that induces a state sequence $\langle s_0, s_1, \dots, s_n \rangle$ such that $s_0 = I$ and, for each i such that $1 \leq i \leq n$, a_i is applicable in s_{i-1} and generates the successor state $s_i = \theta(s_{i-1}, a_i)$. The plan π solves P if and only if $G \subseteq s_n$, i.e. if the goal condition is satisfied following the application of π in I .

We assume that fluents are instantiated from predicates, as in PDDL (Fox and Long 2003). Specifically, there exists

a set of predicates Ψ , each $p \in \Psi$ with an argument list of arity $ar(p)$. Given a set of objects Ω , the set of fluents F is then induced by assigning objects in Ω to the arguments of predicates in Ψ , i.e. $F = \{p(\omega) : p \in \Psi, \omega \in \Omega^{ar(p)}\}$ where, given a set X , X^n is the n -th Cartesian power of X .

Likewise we assume that actions in A are instantiated from operator schema as follows. Let us define $\Omega_v = \{v_1, \dots, v_v\}$ as a new set of objects, $\Omega_v \cap \Omega = \emptyset$, representing variable names. The number of *variable* objects, $|\Omega_v|$, is given by the action with the maximum arity. For instance, in the blockworld $\Omega_v = \{v_1, v_2\}$ since actions *stack* (Figure 1) and *unstack* have two parameters. Let us define a new set of fluents, $F_v \cap F = \emptyset$, that result instantiating Ψ but using only the variable objects in Ω_v .

```
(:action stack
:parameters (?x1 ?x2)
:precondition (and (holding ?x1)
                  (clear ?x2))
:effect (and (not (holding ?x1))
             (not (clear ?x2))
             (clear ?x1)
             (handempty)
             (on ?x1 ?x2)))
```

Figure 1: Example of the *stack* planning action schema from the blockworld as represented in PDDL.

Now we are ready to define an operator schema $\xi \in \Xi$ as:

- The operator *header*: a fluent $header(name(\xi), pars(\xi))$ built using the action name, e.g. $name(\xi) \in \{stack, unstack, pickup, putdown\}$ for the blockworld, and a list of variables, $pars(\xi) \in \Omega_v^{ar(\xi)}$.
- The operator *body* comprises: the *preconditions*, $pre(\xi) \subseteq F_v$, the *positive effects*, $add(\xi) \subseteq F_v$, and the *negative effects*, $del(\xi) \subseteq F_v$ such that, $del(\xi) \subseteq pre(\xi)$, $del(\xi) \cap add(\xi) = \emptyset$ and $pre(\xi) \cap add(\xi) = \emptyset$.

Classical Planning with Conditional Effects

Conditional effects make it possible to repeatedly refer to the same action even though their precise effects depend on the current state. This feature is useful to define our compilation for learning classical planning action models using a classical planner.

In this case each action $a \in A$ has a set of literals $pre(a) \in \mathcal{L}(F)$ called the *precondition* and a set of conditional effects $cond(a)$. Each conditional effect $C \triangleright E \in cond(a)$ is composed of two sets of literals $C \in \mathcal{L}(F)$ (the condition) and $E \in \mathcal{L}(F)$ (the effect). An action $a \in A$ is applicable in state s if and only if $pre(a) \subseteq s$, and the resulting set of *triggered effects* is

$$eff(s, a) = \bigcup_{C \triangleright E \in cond(a), C \subseteq s} E,$$

i.e. effects whose conditions hold in s . The result of applying a in s is a new state $\theta(s, a) = (s \setminus \neg eff(s, a)) \cup eff(s, a)$.

Learning classical planning action models

This section formalizes the different learning tasks addressed in the paper according to the available input knowledge.

On the extreme learning from the complete amount of available input knowledge, the *pre*- and *post*-states of every action execution, is straightforward. In this case the operator schema are derived lifting the literals that change between the pre and post-state of the corresponding action execution. The preconditions are derived lifting the minimal set of literals appearing in all the pre-states that correspond to the same operator.

Learning classical planning action models from plans

This task is formalized as $\Lambda = \langle \Psi, \Pi, \Sigma \rangle$:

- Ψ the set of predicates that define the abstract state space of a given planning domain. This set includes the predicates for describing the actions header.
- $\Pi = \{\pi_1, \dots, \pi_t\}$ is the given set of example plans,
- $\Sigma = \{\sigma_1, \dots, \sigma_t\}$ is a set of labels s.t. each plan π_i , $1 \leq i \leq t$, has an associated label $\sigma_i = (s_i, s'_i)$ where s'_i is the state resulting from executing π_i starting from the state s_i .

A solution to the learning task Λ is a set of operator schema Ξ compliant with the predicates in Ψ , the example plans Π , and their labels Σ .

Learning classical planning action models from states

As explained in this paper we aim to reduce the amount of input knowledge provided to the learning task so we redefine the previous tasks as $\Lambda' = \langle \Psi, \Sigma \rangle$.

A solution to the Λ' learning task is a set of operator schema Ξ that is compliant with the predicates in Ψ , and the given set of initial and final states.

Learning action models from example states using a classical planner

Our approach for addressing $\Lambda' = \langle \Psi, \Sigma \rangle$ is compiling this learning task into a classical planning task P_Λ . The intuition behind the compilation is that a solution to P_Λ is a sequence of actions that first, programs the action model (i.e. the $pre(\xi)$, $del(\xi)$ and $add(\xi)$ for each $\xi \in \Xi$) and then, validates the programmed action model in the given set of labels Σ , one after the other.

To formalize the compilation we first define t classical planning instances $P_1 = \langle F, A, I_1, G_1 \rangle, \dots, P_t = \langle F, A, I_t, G_t \rangle$, that belong to the same planning frame $\Phi = \langle F, A \rangle$ (i.e. share the same fluents and actions and differ only in the initial state and goals). The set of actions is empty $A = \emptyset$ while the set of fluents F is built instantiating the predicates in Ψ with the objects in $\Omega = \{o | o \in s_i \cup s'_i, 1 \leq i \leq t\}$, the set of objects that appear in the states in Σ . Finally the initial state I_i , $1 \leq i \leq t$, is given by the state $s_i \in \sigma_i$ and the goals are defined by the state $s'_i \in \sigma_i$.

Now we are ready to define the compilation for learning action models using an off-the-shelf classical planner. Given a learning task $\Lambda' = \langle \Psi, \Sigma \rangle$ the compilation outputs a classical planning task $P_\Lambda = \langle F_\Lambda, A_\Lambda, I_\Lambda, G_\Lambda \rangle$:

- F_Λ extends F with:
 - Fluents representing the operator header $header(name(\xi), pars(\xi))$ with $pars(\xi) \in \Omega_v^{ar(\xi)}$ for $\xi \in \Xi$.
 - Fluents representing the programmed action model: $pre_f(\xi)$, $del_f(\xi)$ and $add_f(\xi)$ where $f \in F_v$ and $\xi \in \Xi$.
 - Fluents $\{test_i\}_{1 \leq i \leq t}$, indicating the example where the programmed model is currently being validated.
- I_Λ , contains the fluents from F that encode the initial state $s_1 \in P_1$ and the header of the operators we want to learn. For instance, for learning the action model of the blocksworld these headers are $\{pickup(v_1), putdown(v_1), unstack(v_1, v_2), stack(v_1, v_2)\}$.
- $G_\Lambda = \{test_i\}_{1 \leq i \leq t}$, indicates that the programmed action model is validated in all the examples.
- A_Λ replaces the actions in A with actions of three types:
 1. Actions for programming:
 - A *precondition* $f \in F_v$ in the action schema $\xi \in \Xi$:

$$pre(programPre_{f,\xi}) = \{\neg pre_f(\xi), \neg del_f(\xi), \neg add_f(\xi)\},$$

$$cond(programPre_{f,\xi}) = \{\emptyset\} \triangleright \{pre_f(\xi)\}.$$
 - A *negative effect* $f \in F_v$ in the action schema $\xi \in \Xi$:

$$pre(programDel_{f,\xi}) = \{pre_f(\xi), \neg del_f(\xi), \neg add_f(\xi)\},$$

$$cond(programDel_{f,\xi}) = \{\emptyset\} \triangleright \{del_f(\xi)\}.$$
 - A *positive effect* $f \in F_v$ in the action schema $\xi \in \Xi$:

$$pre(programAdd_{f,\xi}) = \{\neg pre_f(\xi), \neg del_f(\xi), \neg add_f(\xi)\},$$

$$cond(programAdd_{f,\xi}) = \{\emptyset\} \triangleright \{add_f(\xi)\}.$$
 2. Actions for applying an already programmed operator schema $\xi \in \Xi$ bounding it with objects $\omega \subseteq \Omega$

$$pre(apply_{\xi,\omega}) = \{header(name(\xi), pars(\xi))\} \cup$$

$$\{pre_f(\xi) \implies p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))},$$

$$cond(apply_{\xi,v,v'}) = \{del_f(\xi)\} \triangleright \{\neg p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))},$$

$$\{add_f(\xi)\} \triangleright \{p(\omega)\}_{\forall p \in \Psi, f=p(pars(\xi))}.$$
 3. Actions for changing the active example where the action model is currently being validated.

$$pre(validate_i) = G_i \cup \{test_j\}_{j \in [1 \leq j < i]},$$

$$cond(validate_i) = \{\emptyset\} \triangleright \{test_i\}.$$

Lemma 1. Any classical plan π that solves P_Λ induces a valid action model that solves the learning task Λ .

Proof sketch. Once an operator schema is programmed it cannot be modified and can only be applied. The only way of achieving a test fluent is by applying a sequence of programmed operator schema until achieving the goal state in its associated label starting from the initial state of the corresponding label. If this is done for all the input examples it means that the programmed action model is compliant with the learning input knowledge and hence, that solves the action model learning task. \square

Interestingly this compilation is valid for partially specified action models since known preconditions and effects (fluents $pre_f(\xi)$, $del_f(\xi)$ and $add_f(\xi)$) can be part of the initial state I_Λ . With this regard the approach allows also transfer learning in which we generate the action model for a given sub-task and then encode this model as already programmed actions for learning new or more challenging action models.

Learning action models from example plans using a classical planner

The compilation can be extended to the learning scenario defined by Λ in which a set of action plans is also available. In this scenario each plan $\pi_i \in \Pi$, $1 \leq i \leq t$, is a solution to the corresponding classical planning instance $P_i = \langle F, A, I_i, G_i \rangle$. The compilation extensions are:

- F_Λ includes the set of fluents $F_\Pi = \{plan(name(\xi), j, \Omega^{ar(\xi)})\}$ for encoding the j steps of the $1 \leq i \leq t$ plans in Π with $F_{\Pi_i} \subseteq F_\Pi$ the subset of fluents encoding the plan corresponding to the i^{th} example. In addition fluents at_j and $next_{j,j_2}$, $1 \leq j < j_2 \leq n$, indicate the plan step where the programmed action model is being validated and n is the max length of a plan in Π .
- I_Λ is extended with the fluents from F_{Π_1} that encode the plan $\pi_1 \in \Pi$ for solving P_1 , and the fluents at_1 and $\{next_{j,j_2}\}$, $1 \leq j < j_2 \leq n$, for indicating the plan step where to start validating the programmed action model. Goals are like in the original compilation.
- With respect to the actions,
 1. The actions for programming the preconditions/effects of a given operator schema are the same.
 2. The actions for applying an operator schema have an extra precondition $f \in F_{\Pi_i}$ s.t. f encodes the current plan step and extra conditional effect $\{at_j\} \triangleright \{\neg at_j, at_{j+1}\}_{\forall j \in [1,n]}$ for advancing the current plan step.
 3. The actions for changing the active test have an extra precondition, $at_{|\Pi_i|}$, to indicate that we simulated the full current plan and extra conditional effects to load the next plan where to validate the programmed action model:

$$\{f\} \triangleright \{\neg f\}_{f \in F_{\Pi_i}},$$

$$\{\emptyset\} \triangleright \{f\}_{f \in F_{\Pi_i+1}},$$

$$\{\emptyset\} \triangleright \{\neg at_{|\Pi_i|}, at_1\}.$$

Generating informative learning examples

TBD.

Observation: If a predicate is not appearing either in the initial not in the final state it will not appear in the learned action model.

Question: Can we propose a search algorithm that generates states with the full diversity of predicates of a given domain?

Evaluation

Learning action models from example plans

The performance of our learning approach is evaluated for different levels of available input knowledge and using different sources for collecting this initial knowledge. In all the cases we assess the performance of our learning approach using the cardinality of the *symmetric difference* sets that are computed between the set of preconditions, del and add effects (1), in the learned model and (2), in the actual models. In all the experiments the compilation is solved using the SAT-based classical planner MADAGASCAR (Rintanen 2014).

Table 1 shows the mean error and standard deviation of the learned models with respect to the actual action models when (1) using *hand-picked* examples, (2) examples collected using the classical planner FAST-DOWNWARD (Helmert 2006) and (3) examples collected *randomly*. The standard deviation provides a measure of how this error is distributed among the different operators in the domain. If this deviation is 0 it means that is equally distributed in all the domain operators.

Learning action models from example states

TBD.

Related work

Conclusions

This paper presents a novel approach for learning classical planning action models from minimal input knowledge and using exclusively existing classical planners. Using classical planning focus the learning task in the more succinct action models given that classical planners have a preference for short plans.

Learning action models from examples allows the reformulation of a domain theory. An interesting research direction is the study of domain reformulation using features that allow more compact solutions like the *reachable* or *movable* features in the Sokoban domain.

References

- Cresswell, S. N.; McCluskey, T. L.; and West, M. M. 2013. Acquiring planning domain models using locm. *The Knowledge Engineering Review* 28(02):195–213.
- Fern, A.; Yoon, S. W.; and Givan, R. 2004. Learning domain-specific control knowledge from random walks. In *ICAPS*, 191–199.
- Fox, M., and Long, D. 2003. Pddl2. 1: An extension to pddl for expressing temporal planning domains. *J. Artif. Intell. Res.(JAIR)* 20:61–124.
- Francs, G.; Ramrez, M.; Lipovetzky, N.; and Geffner, H. 2017. Purely declarative action representations are overrated: Classical planning with simulators. In *International Joint Conference on Artificial Intelligence*.
- Geffner, H., and Bonet, B. 2013. A concise introduction to models and methods for automated planning.

- Helmert, M. 2006. The fast downward planning system. *J. Artif. Intell. Res.(JAIR)* 26:191–246.
- Kambhampati, S. 2007. Model-lite planning for the web age masses: The challenges of planning with incomplete and evolving domain models. In *Proceedings of the National Conference on Artificial Intelligence*.
- Michalski, R. S.; Carbonell, J. G.; and Mitchell, T. M. 2013. *Machine learning: An artificial intelligence approach*. Springer Science & Business Media.
- Rintanen, J. 2014. Madagascar: Scalable planning with sat. *Proceedings of the 8th International Planning Competition (IPC-2014)*.
- Segovia-Aguas, J.; Jiménez, S.; and Jonsson, A. 2017. Generating context-free grammars using classical planning. In *International Joint Conference on Artificial Intelligence*.
- Yang, Q.; Wu, K.; and Jiang, Y. 2007. Learning action models from plan examples using weighted max-sat. *Artificial Intelligence* 171(2-3):107–143.

	Preconditions				Del Effects				Add Effects		
	Hand	Planner	Rand		Hand	Planner	Rand		Hand	Planner	Rand
Blocks	5.00 ± 8.66	58.33 ± 43.30			0.00	16.70 ± 3.36			6.07 ± 6.26	19.38 ± 19.40	
Gripper	55.56 ± 7.86	40.00 ± 20.00			0.00	9.09			35.98 ± 26.38	72.73	
Miconic	37.62 ± 16.52	61.67 ± 2.89			0.00	8.33 ± 14.43			61.36 ± 13.30	29.17 ± 18.16	
Visitall	50.00	100.00			0.00	20.00			0.00	0.00	

Table 1: Mean error and standard deviation of the learned models when using hand-picked examples and examples collected using the classical planner Fast-Downward.